

LSC - lab 6 - Kubernetes

Author: Michał Skalka

GitHub repository: <https://github.com/Skalakid/large-scale-computing>

Run the lab using run.sh script: ./run.sh

```
config.sh U X
large-scale-computing > lab6 > config.sh
1  brew install minikube
2  brew install kind
3
4  minikube start
5
6  kubectl version --client kubectl get nodes
7
8  curl https://raw.githubusercontent.com/helm/helm/master/scripts/get-helm-3 | bash
9  helm repo add stable https://charts.helm.sh/stable
10 helm repo update
11
12 helm install nfs-server stable/nfs-server-provisioner \
13   --set persistence.enabled=true \
14   --set persistence.size=1Gi \
15   --set storageClass.name=nfs-storage
16
17 kubectl get storageclass
18
19
20 kubectl apply -f pvc.yaml
21 kubectl apply -f nginx-deployment.yaml
22 kubectl apply -f nginx-service.yaml
23 kubectl apply -f content-job.yaml
24
25 minikube service nginx-service
```

1. Create a k8s cluster using Minikube

```
minikube start
minikube v1.35.0 on Darwin 15.3 (arm64)
Automatically selected the docker driver
Using Docker Desktop driver with root privileges
Starting "minikube" primary control-plane node in "minikube" cluster
Pulling base image v0.0.46 ...
Downloading Kubernetes v1.32.0 preload ...
> preloaded-images-k8s-v18-v1...: 314.92 MiB / 314.92 MiB 100.00% 5.01 Mi
> gcr.io/k8s-minikube/kicbase...: 452.84 MiB / 452.84 MiB 100.00% 5.29 Mi
Creating docker container (CPUs=2, Memory=4600MB) ...
Preparing Kubernetes v1.32.0 on Docker 27.4.1 ...
  ▪ Generating certificates and keys ...
  ▪ Booting up control plane ...
  ▪ Configuring RBAC rules ...
Configuring bridge CNI (Container Networking Interface) ...
Verifying Kubernetes components...
  ▪ Using image gcr.io/k8s-minikube/storage-provisioner:v5
Enabled addons: default-storageclass, storage-provisioner
Done! kubectl is now configured to use "minikube" cluster and "default" namespace by default
```

2. Using Helm, install an NFS server and provisioner in the cluster.

The NFS (Network File System) server provides shared storage that can be accessed by multiple pods. We installed it in the cluster using a Helm chart called `nfs-server-provisioner`. It dynamically provisions Persistent Volumes (PVs) on demand. The Persistent Volume is a storage resource in the cluster provided by the NFS server. It abstracts the actual storage (NFS in this case) from the Kubernetes applications.

```
~/Desktop/Studia/SEM_8/LCS/lab6 ➤ curl https://raw.githubusercontent.com/helm/helm/master/scripts/get-helm-3 | bash
% Total    % Received % Xferd  Average Speed   Time    Time     Time  Current
                                 Dload  Upload   Total   Spent    Left   Speed
100 11913  100 11913    0     0  74823      0  --:--:-- --:--:-- --:--:--  74456
Downloading https://get.helm.sh/helm-v3.17.3-darwin-arm64.tar.gz
Verifying checksum... Done.
Preparing to install helm into /usr/local/bin
Password:
helm installed into /usr/local/bin/helm
~/Desktop/Studia/SEM_8/LCS/lab6 ➤ helm repo add stable https://charts.helm.sh/stable
"stable" has been added to your repositories
~/Desktop/Studia/SEM_8/LCS/lab6 ➤ helm repo update
Hang tight while we grab the latest from your chart repositories...
...Successfully got an update from the "stable" chart repository
Update Complete. ✨Happy Helming!✨
~/Desktop/Studia/SEM_8/LCS/lab6 ➤ helm install nfs-server stable/nfs-server-provisioner \
--set persistence.enabled=true \
--set persistence.size=1Gi \
--set storageClass.name=nfs-storage
WARNING: This chart is deprecated
NAME: nfs-server
LAST DEPLOYED: Tue Apr 15 20:16:12 2025
NAMESPACE: default
STATUS: deployed
REVISION: 1
TEST SUITE: None
NOTES:
The NFS Provisioner service has now been installed.

A storage class named 'nfs-storage' has now been created
and is available to provision dynamic volumes.

You can use this storageclass by creating a `PersistentVolumeClaim` with the
correct storageClassName attribute. For example:

---
kind: PersistentVolumeClaim
apiVersion: v1
metadata:
  name: test-dynamic-volume-claim
spec:
  storageClassName: "nfs-storage"
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 100Mi
```

```
~/Desktop/Studia/SEM_8/LCS/lab6 ➤ kubectl get storageclass
NAME                PROVISIONER                                RECLAIMPOLICY  VOLUMEBINDINGMODE  ALLOWVOLUMEEXPANSION  AGE
nfs-storage         cluster.local/nfs-server-nfs-server-provisioner  Delete         Immediate           true                  94s
standard (default)  k8s.io/minikube-hostpath                    Delete         Immediate           false                 3m50s
```

3. Create a Persistent Volume Claim which will bind to a NFS Persistent Volume provisioned dynamically by the provisioner installed in the previous step

Persistent Volume Claim (PVC) is a request for storage by a user/application. It will dynamically bind to the NFS Provisioner provisioned in the previous step.

```
pvc.yaml ×
pvc.yaml > {} spec > {} resources > {} requests
1  apiVersion: v1
2  kind: PersistentVolumeClaim
3  metadata:
4    name: nfs-pvc
5  spec:
6    accessModes:
7      - ReadWriteMany
8    storageClassName: nfs-storage
9    resources:
10     requests:
11     storage: 500Mi
12
```

```
~/Desktop/Studia/SEM_8/LCS/lab6 ➤ kubectl apply -f pvc.yaml
persistentvolumeclaim/nfs-pvc created
```

4. Create a Deployment with a HTTP server (e.g., apache or nginx). The web content directory should be mounted as a volume using the PVC created in the previous step

The Deployment manages the lifecycle of a set of identical pods (like the HTTP server). In this case, it deploys an nginx web server and ensures that the pod is running at all times. The deployment uses the PVC to mount the shared volume inside the pod.

```
pvc.yaml nginx-deployment.yaml X
nginx-deployment.yaml > {} spec > {} template > {} spec > [ ] vol
1  apiVersion: apps/v1
2  kind: Deployment
3  metadata:
4    name: nginx-deployment
5  spec:
6    replicas: 1
7    selector:
8      matchLabels:
9        app: nginx
10   template:
11     metadata:
12       labels:
13         app: nginx
14     spec:
15       containers:
16         - name: nginx
17           image: nginx
18           ports:
19             - containerPort: 80
20           volumeMounts:
21             - name: web-content
22               mountPath: /usr/share/nginx/html
23       volumes:
24         - name: web-content
25           persistentVolumeClaim:
26             claimName: nfs-pvc
27
```

```
~/Desktop/Studia/SEM_8/LCS/lab6 > kubectl apply -f nginx-deployment.yaml
deployment.apps/nginx-deployment created
~/Desktop/Studia/SEM_8/LCS/lab6 > █
```

5. Create a Service associated with the Pod(s) of the HTTP server Deployment.

A Service in Kubernetes is an abstraction that exposes a set of pods as a network service. In this case, the service is configured to expose the HTTP server pod(s) created by the Deployment

```
pvc.yaml  nginx-deployment.yaml  nginx-service.yaml X
nginx-service.yaml > {} spec > [ ] ports > {} 0
1  apiVersion: v1
2  kind: Service
3  metadata:
4    name: nginx-service
5  spec:
6    type: NodePort
7    selector:
8      app: nginx
9    ports:
10     - port: 80
11       targetPort: 80
12       nodePort: 30080
13
```

```
~/Desktop/Studia/SEM_8/LCS/lab6  kubectl apply -f nginx-service.yaml
service/nginx-service created
~/Desktop/Studia/SEM_8/LCS/lab6  minikube service nginx-service
|-----|
| NAMESPACE | NAME       | TARGET PORT | URL                               |
|-----|
| default   | nginx-service | 80          | http://192.168.49.2:30080        |
|-----|
Starting tunnel for service nginx-service.
|-----|
| NAMESPACE | NAME       | TARGET PORT | URL                               |
|-----|
| default   | nginx-service |             | http://127.0.0.1:65472          |
|-----|
Opening service default/nginx-service in default browser...
! Because you are using a Docker driver on darwin, the terminal needs to be open to run it.
```

6. Create a Job which mounts the PVC and copies a sample content through the shared NFS PV.

A Job in Kubernetes is used to run a task to completion. In this case, the Job is responsible for mounting the same PVC used by the HTTP server and copying sample index.htm into it.

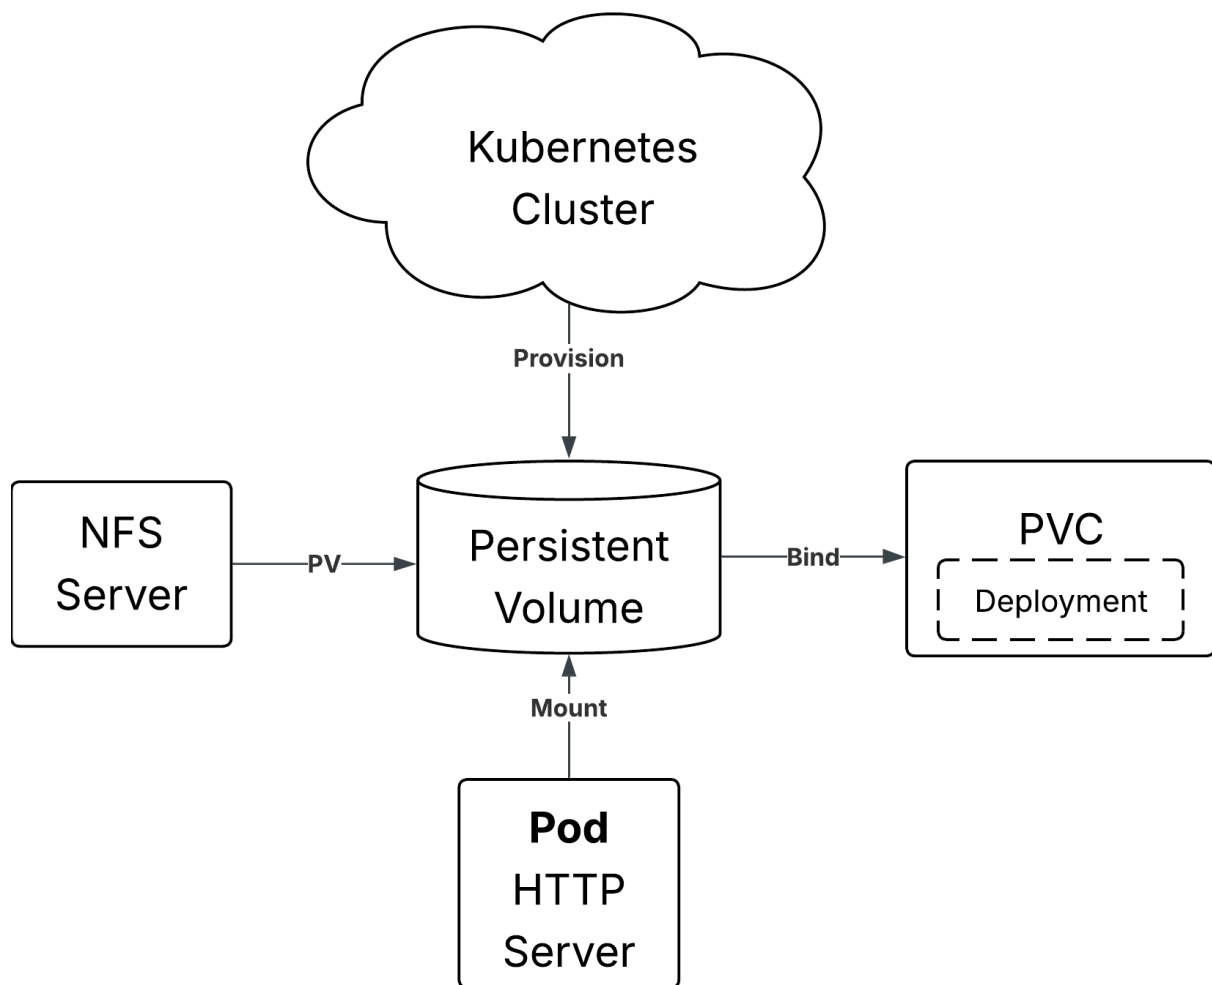
```
pvc.yaml nginx-deployment.yaml nginx-service.yaml content-job.yaml X
content-job.yaml > {} spec
1  apiVersion: batch/v1
2  kind: Job
3  metadata:
4    name: copy-content-job
5  spec:
6    template:
7      spec:
8        containers:
9          - name: copy
10           image: busybox
11           command:
12             [
13               "sh",
14               "-c",
15               "echo '<h1>Hello from NFS!</h1>' > /mnt/data/index.html",
16             ]
17           volumeMounts:
18             - name: web-content
19               mountPath: /mnt/data
20         volumes:
21           - name: web-content
22             persistentVolumeClaim:
23               claimName: nfs-pvc
24             restartPolicy: Never
25         backoffLimit: 2
26
```

```
~/Desktop/Studia/SEM_8/LCS/lab6 kubectl apply -f content-job.yaml
job.batch/copy-content-job created
~/Desktop/Studia/SEM_8/LCS/lab6 minikube service nginx-service
-----
| NAMESPACE | NAME       | TARGET PORT | URL                |
|-----|-----|-----|-----|
| default   | nginx-service | 80          | http://192.168.49.2:30080 |
|-----|-----|-----|-----|
Starting tunnel for service nginx-service.
-----
| NAMESPACE | NAME       | TARGET PORT | URL                |
|-----|-----|-----|-----|
| default   | nginx-service |            | http://127.0.0.1:65524 |
|-----|-----|-----|-----|
Opening service default/nginx-service in default browser...
! Because you are using a Docker driver on darwin, the terminal needs to be open to run it.
```

i http://127.0.0.1:65524

Hello from NFS!

Diagram:



Provision: NFS provisioner automatically creates a Persistent Volume when a PVC is detected.

Bind: Kubernetes matches and binds the PVC to an available PV.

Mount: The PV is mounted into the pod so the container can use it as a normal directory.